

Ensemble-Techniques-Project

November 11, 2019

1 Ensemble Techniques Project

1.1 Data Description

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed.

1.2 Domain

Banking

1.3 Context

Leveraging customer information is paramount for most businesses. In the case of a bank, attributes of customers like the ones mentioned below can be crucial in strategizing a marketing campaign when launching a new product.

1.4 Attribute Information

1. age (numeric)
2. job : type of job (categorical: 'admin.', 'bluecollar', 'entrepreneur', 'housemaid', 'management', 'retired', 'selfemp', 'services', 'technician', 'unemployed', 'unknown')
3. marital : marital status (categorical: 'divorced', 'married', 'single', 'unknown'; note: 'divorced' means divorced or widowed)
4. education (categorical: 'basic.4y', 'basic.6y', 'basic.9y', 'high.school', 'illiterate', 'professional.course', 'university.degree', 'unknown')
5. default: has credit in default? (categorical: 'no', 'yes', 'unknown')
6. balance: average yearly balance, in euros (numeric)
7. housing: has housing loan? (categorical: 'no', 'yes', 'unknown')
8. loan: has personal loan? (categorical: 'no', 'yes', 'unknown')
9. contact: contact communication type (categorical: 'cellular', 'telephone')
10. day: last contact day of the month (numeric 1 -31)
11. month: last contact month of year (categorical: 'jan', 'feb', 'mar', ..., 'nov', 'dec')
12. duration: last contact duration, in seconds (numeric). Important note: this attribute highly affects the output target (e.g., if duration=0 then y='no'). Yet, the duration is not known before a call is performed. Also, after the end of the call y is obviously known. Thus, this input

should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model.

13. campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
14. pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted)
15. previous: number of contacts performed before this campaign and for this client (numeric)
16. poutcome: outcome of the previous marketing campaign (categorical: 'failure', 'nonexistent', 'success')
17. target: has the client subscribed a term deposit? (binary: "yes", "no")

1.5 Learning Outcomes

- Exploratory Data Analysis
- Preparing the data to train a model
- Training and making predictions using an Ensemble Model
- Tuning an Ensemble model

1.6 Objective

The classification goal is to predict if the client will subscribe (yes/no) a term deposit (variable y).

1.6.1 Step 1: Import the necessary libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

[2]: # suppress warnings
import warnings
warnings.filterwarnings("ignore")
```

1.6.2 Read the data as a data frame

```
[3]: !ls

bank-full.csv          'Problem statement- ensemble project.pdf'
Ensemble-Techniques-Project.ipynb

[4]: bank = pd.read_csv("bank-full.csv")
[5]: bank.head()
[5]:  age      job  marital  education  default  balance  housing  loan  \
0   58  management  married   tertiary     no     2143     yes    no
1   44  technician  single    secondary     no      29     yes    no
2   33  entrepreneur  married    secondary     no      2     yes    yes
```

3	47	blue-collar	married	unknown	no	1506	yes	no
4	33	unknown	single	unknown	no	1	no	no

	contact	day	month	duration	campaign	pdays	previous	poutcome	Target
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no

1.6.3 Perform basic EDA which should include the following and print out your insights at every step

Shape of the data

```
[6]: print("Shape of the dataframe: {}".format(bank.shape))
```

Shape of the dataframe: (45211, 17)

The dataframe has 45211 samples and 17 features for each sample.

Data type of each attribute

```
[7]: bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
age                45211 non-null int64
job                45211 non-null object
marital            45211 non-null object
education          45211 non-null object
default            45211 non-null object
balance            45211 non-null int64
housing            45211 non-null object
loan               45211 non-null object
contact            45211 non-null object
day                45211 non-null int64
month              45211 non-null object
duration           45211 non-null int64
campaign           45211 non-null int64
pdays             45211 non-null int64
previous           45211 non-null int64
poutcome           45211 non-null object
Target             45211 non-null object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

Note that there are some categorical variables like job, marital, etc. that are currently in object format that need to be mapped to 0 and 1 or using one hot embedding.

Checking the presence of missing values

```
[8]: # Print the percentage of missing values
bank.isna().sum()/len(bank.index) * 100
```

```
[8]: age          0.0
     job          0.0
     marital      0.0
     education    0.0
     default      0.0
     balance      0.0
     housing      0.0
     loan         0.0
     contact      0.0
     day          0.0
     month        0.0
     duration     0.0
     campaign     0.0
     pdays        0.0
     previous     0.0
     poutcome     0.0
     Target       0.0
     dtype: float64
```

Note that there is no missing value in any of the rows. But, as we can see from the attribute information above, there are values like “unknown”, “other”, etc. which serve as missing values.

Let’s start off by printing the unique values for each categorical attribute.

```
[9]: bank.job.value_counts()
```

```
[9]: blue-collar    9732
     management   9458
     technician   7597
     admin.       5171
     services     4154
     retired      2264
     self-employed 1579
     entrepreneur 1487
     unemployed   1303
     housemaid    1240
     student      938
     unknown      288
     Name: job, dtype: int64
```

There are 281 rows with unknown jobs. Let’s remove such rows.

```
[10]: bank.drop(bank[bank.job=="unknown"].index,axis=0,inplace=True)
```

```
[11]: bank.marital.value_counts()
```

```
[11]: married    27011
     single    12722
```

```
divorced      5190
Name: marital, dtype: int64
```

```
[12]: bank.education.value_counts()
```

```
[12]: secondary      23131
      tertiary      13262
      primary       6800
      unknown       1730
      Name: education, dtype: int64
```

There are 1822 rows with missing education data. Let's remove these rows.

```
[13]: bank.drop(bank[bank.education=="unknown"].index,axis=0,inplace=True)
```

```
[14]: bank.default.value_counts()
```

```
[14]: no      42411
      yes      782
      Name: default, dtype: int64
```

```
[15]: bank.contact.value_counts()
```

```
[15]: cellular      28213
      unknown      12286
      telephone    2694
      Name: contact, dtype: int64
```

Note the high number of unknown entries in contact. It's better to drop the entire column.

```
[16]: bank.drop("contact",axis=1,inplace=True)
```

```
[17]: bank.month.value_counts()
```

```
[17]: may      13192
      jul      6601
      aug      6037
      jun      4980
      nov      3842
      apr      2820
      feb      2533
      jan      1318
      oct       690
      sep       532
      mar       448
      dec       200
      Name: month, dtype: int64
```

```
[18]: bank.poutcome.value_counts()
```

```
[18]: unknown      35286
      failure      4709
      other        1774
      success      1424
```

Name: poutcome, dtype: int64

Again, note the high number of unknown entries in poutcome attribute. It's better to drop the entire column.

```
[19]: bank.drop("poutcome",axis=1,inplace=True)
```

5 Point summary of numerical attributes

```
[20]: bank.describe()
```

```
[20]:
```

	age	balance	day	duration	campaign \
count	43193.000000	43193.000000	43193.000000	43193.000000	43193.000000
mean	40.764082	1354.027342	15.809414	258.323409	2.758178
std	10.512640	3042.103625	8.305970	258.162006	3.063987
min	18.000000	-8019.000000	1.000000	0.000000	1.000000
25%	33.000000	71.000000	8.000000	103.000000	1.000000
50%	39.000000	442.000000	16.000000	180.000000	2.000000
75%	48.000000	1412.000000	21.000000	318.000000	3.000000
max	95.000000	102127.000000	31.000000	4918.000000	58.000000

	pdays	previous
count	43193.000000	43193.000000
mean	40.404070	0.584863
std	100.420624	2.332672
min	-1.000000	0.000000
25%	-1.000000	0.000000
50%	-1.000000	0.000000
75%	-1.000000	0.000000
max	871.000000	275.000000


Let's start off by noting the presence of outliers in numerical attributes.

1. balance column has a minimum value of -8019 which is negative, but the average annual balance should not be negative. This either can be a typo or an outlier.
2. There is a significant difference in 75% and max values of age attribute.
3. Similarly, there are outliers in duration, campaign.
4. Note that previous and pdays have majority of entries with -1 or 0 (missing data or client not contacted), we can drop this column.

Checking the presence of outliers I will use box plot to check the presence of outliers.

Age

```
[21]: plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
sns.distplot(bank["age"])
plt.subplot(1,2,2)
sns.boxplot(bank["age"])
plt.show()
```



```
[22]: bank.age.max()
```

```
[22]: 95
```

```
[23]: bank.age.min()
```


```
[23]: 18
```

While there are outliers in this attribute, but both max and min ages (95,18) are realistic, and thus, we won't remove them.

Balance

We can use zscore as a standard for outlier removal. Any value which lies 3σ away from the mean can be considered an outlier and dropped.

```
[24]: plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
sns.distplot(bank["balance"])
plt.subplot(1,2,2)
sns.boxplot(bank["balance"])
plt.show()
```



```
[25]: from scipy.stats import zscore
print("Current mean: {}".format(bank.balance.mean()))
```

Current mean: 1354.0273423934434

```
[26]: bank["balance_outliers"] = zscore(bank["balance"])
```

```
[27]: bank.drop(bank[(bank["balance_outliers"]>3) | (bank["balance_outliers"] < -3)].
→index,axis=0,inplace=True)
```

```
[28]: bank.shape
```


```
[28]: (42484, 16)
```

```
[29]: print("Rows dropped: {}".format(45211-44466))
```

Rows dropped: 745

```
[30]: # We don't need the zscore column anymore
bank.drop("balance_outliers",axis=1,inplace=True)
```

```
[31]: plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
sns.distplot(bank["balance"])
plt.subplot(1,2,2)
sns.boxplot(bank["balance"])
plt.show()
```


We can clearly see a few more outliers to the negative side. Let's drop them.

```
[32]: sum(bank.balance<-2500)
```

```
[32]: 8
```



```
[33]: bank.drop(bank[bank.balance<-2500].index,axis=0,inplace=True)
```

```
[34]: plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
sns.distplot(bank["balance"])
plt.subplot(1,2,2)
sns.boxplot(bank["balance"])
plt.show()
```



Day


```
[35]: plt.figure(figsize=(12,6))  
plt.subplot(1,2,1)  
sns.distplot(bank["day"])  
plt.subplot(1,2,2)  
sns.boxplot(bank["day"])  
plt.show()
```



The day attribute does not have any outliers.

Duration

```
[36]: plt.figure(figsize=(12,6))  
plt.subplot(1,2,1)  
sns.distplot(bank["duration"])  
plt.subplot(1,2,2)  
sns.boxplot(bank["duration"])  
plt.show()
```




We know that duration is not known in advance and by the end of the call, when we know the duration, we also know the result. So, we can drop this attribute. That's why, we are not concerned about the outlier presence in duration attribute.

```
[37]: bank.drop("duration",axis=1,inplace=True)
```


Pdays

```
[38]: plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
sns.distplot(bank["pdays"])
plt.subplot(1,2,2)
sns.boxplot(bank["pdays"])
plt.show()
```



Because of the unbalanced distribution in value ranges for pdays attribute, it's better to convert it to categorical variable by binning the values.

```
[39]: plt.figure(figsize=(12,6))  
plt.subplot(1,2,1)  
sns.distplot(bank[bank["pdays"]<400].pdays)  
plt.subplot(1,2,2)  
sns.boxplot(bank[bank["pdays"]<400].pdays)  
plt.show()
```




Note the large number of entries for pdays having a specific value. This makes it relevant for dropping.

```
[40]: bank.drop("pdays",axis=1,inplace=True)
```

Previous

```
[41]: plt.figure(figsize=(12,6))  
plt.subplot(1,2,1)  
sns.distplot(bank["previous"])  
plt.subplot(1,2,2)  
sns.boxplot(bank["previous"])  
plt.show()
```




```
[42]: sum(bank.previous>30)
```

```
[42]: 12
```

Let's drop all the rows with values more than 30 for previous.

```
[43]: bank.drop(bank[bank.previous>30].index,axis=0,inplace=True)
```

```
[44]: plt.figure(figsize=(12,6))
plt.subplot(1,2,1)
sns.distplot(bank["previous"])
plt.subplot(1,2,2)
sns.boxplot(bank["previous"])
plt.show()
```



```
[45]: sum(bank.previous==0)
```

```
[45]: 34703
```

Note that a large number of rows have previous as 0. It's better to drop this attribute for now.

```
[46]: bank.drop("previous",axis=1,inplace=True)
```

```
[47]: bank.shape
```

```
[47]: (42464, 12)
```

```
[48]: print("Rows deleted: {}".format(45211-bank.shape[0]))
```

```
Rows deleted: 2747
```

1.6.4 Data Preparation

Since we have already dealt with missing data and the data types, let's directly focus on embedding the data.

```
[49]: bank.head()
```

```
[49]:   age      job  marital  education  default  balance  housing  loan  day  \
0   58  management  married   tertiary     no    2143     yes    no    5
1   44  technician  single   secondary     no     29     yes    no    5
2   33  entrepreneur  married   secondary     no      2     yes   yes    5
5   35  management  married   tertiary     no    231     yes    no    5
6   28  management  single   tertiary     no    447     yes   yes    5
```

```
   month  campaign  Target
0   may          1     no
1   may          1     no
2   may          1     no
5   may          1     no
6   may          1     no
```

```
[50]: bank.job.unique()
```

```
[50]: array(['management', 'technician', 'entrepreneur', 'retired', 'admin.',
        'services', 'blue-collar', 'self-employed', 'unemployed',
        'housemaid', 'student'], dtype=object)
```

```
[51]: # Perform one hot embedding and remove original job column
bank = pd.concat([bank,
                  pd.get_dummies(bank.job,drop_first=True)
                  ], axis=1).drop("job",axis=1)
bank.head()
```

```
[51]:   age  marital  education  default  balance  housing  loan  day  month  campaign  \
0   58  married   tertiary     no    2143     yes    no    5   may          1
1   44   single   secondary     no     29     yes    no    5   may          1
```

2	33	married	secondary	no	2	yes	yes	5	may	1
5	35	married	tertiary	no	231	yes	no	5	may	1
6	28	single	tertiary	no	447	yes	yes	5	may	1

		blue-collar	entrepreneur	housemaid	management	retired	\
0	...	0	0	0	1	0	
1	...	0	0	0	0	0	
2	...	0	1	0	0	0	
5	...	0	0	0	1	0	
6	...	0	0	0	1	0	

		self-employed	services	student	technician	unemployed
0		0	0	0	0	0
1		0	0	0	1	0
2		0	0	0	0	0
5		0	0	0	0	0
6		0	0	0	0	0

[5 rows x 21 columns]

```
[52]: bank.marital.unique()
```

```
[52]: array(['married', 'single', 'divorced'], dtype=object)
```

```
[53]: # Perform one hot embedding and remove original marital column
bank = pd.concat([bank,
                  pd.get_dummies(bank.marital,drop_first=True)
                  ], axis=1).drop("marital",axis=1)
bank.head()
```

	age	education	default	balance	housing	loan	day	month	campaign	Target	\
0	58	tertiary	no	2143	yes	no	5	may	1	no	
1	44	secondary	no	29	yes	no	5	may	1	no	
2	33	secondary	no	2	yes	yes	5	may	1	no	
5	35	tertiary	no	231	yes	no	5	may	1	no	
6	28	tertiary	no	447	yes	yes	5	may	1	no	

		housemaid	management	retired	self-employed	services	student	\
0	...	0	1	0	0	0	0	
1	...	0	0	0	0	0	0	
2	...	0	0	0	0	0	0	
5	...	0	1	0	0	0	0	
6	...	0	1	0	0	0	0	

		technician	unemployed	married	single
0		0	0	1	0
1		1	0	0	1
2		0	0	1	0
5		0	0	1	0

```
6          0          0          0          1
```

```
[5 rows x 22 columns]
```

```
[54]: bank.education.unique()
```

```
[54]: array(['tertiary', 'secondary', 'primary'], dtype=object)
```

```
[55]: # Perform one hot embedding and remove original education column
bank = pd.concat([bank,
                  pd.get_dummies(bank.education,drop_first=True)
                  ], axis=1).drop("education",axis=1)
bank.head()
```

```
[55]:   age default  balance housing loan  day month  campaign Target  blue-collar \
0   58      no    2143     yes   no    5   may         1     no         0
1   44      no      29     yes   no    5   may         1     no         0
2   33      no       2     yes  yes    5   may         1     no         0
5   35      no    231     yes   no    5   may         1     no         0
6   28      no    447     yes  yes    5   may         1     no         0

   ...  retired  self-employed  services  student  technician  unemployed  \
0   ...        0              0         0         0           0           0
1   ...        0              0         0         0           1           0
2   ...        0              0         0         0           0           0
5   ...        0              0         0         0           0           0
6   ...        0              0         0         0           0           0

   married  single  secondary  tertiary
0         1       0         0         1
1         0       1         1         0
2         1       0         1         0
5         1       0         0         1
6         0       1         0         1
```

```
[5 rows x 23 columns]
```

```
[56]: bank.default.unique()
```

```
[56]: array(['no', 'yes'], dtype=object)
```

```
[57]: bank.default = bank.default.map({"yes":1,"no":0})
```

```
[58]: bank.head()
```

```
[58]:   age  default  balance housing loan  day month  campaign Target  \
0   58         0    2143     yes   no    5   may         1     no
1   44         0      29     yes   no    5   may         1     no
2   33         0       2     yes  yes    5   may         1     no
5   35         0    231     yes   no    5   may         1     no
6   28         0    447     yes  yes    5   may         1     no
```


	blue-collar	...	retired	self-employed	services	student	technician	\
0	0	...	0	0	0	0	0	
1	0	...	0	0	0	0	1	
2	0	...	0	0	0	0	0	
5	0	...	0	0	0	0	0	
6	0	...	0	0	0	0	0	

	unemployed	married	single	secondary	tertiary
0	0	1	0	0	1
1	0	0	1	1	0
2	0	1	0	1	0
5	0	1	0	0	1
6	0	0	1	0	1

[5 rows x 23 columns]

```
[59]: bank.housing.unique()
```

```
[59]: array(['yes', 'no'], dtype=object)
```

```
[60]: bank.housing = bank.housing.map({"yes":1,"no":0})
```

```
[61]: bank.head()
```

```
[61]:
```

	age	default	balance	housing	loan	day	month	campaign	Target	\
0	58	0	2143	1	no	5	may	1	no	
1	44	0	29	1	no	5	may	1	no	
2	33	0	2	1	yes	5	may	1	no	
5	35	0	231	1	no	5	may	1	no	
6	28	0	447	1	yes	5	may	1	no	

	blue-collar	...	retired	self-employed	services	student	technician	\
0	0	...	0	0	0	0	0	
1	0	...	0	0	0	0	1	
2	0	...	0	0	0	0	0	
5	0	...	0	0	0	0	0	
6	0	...	0	0	0	0	0	

	unemployed	married	single	secondary	tertiary
0	0	1	0	0	1
1	0	0	1	1	0
2	0	1	0	1	0
5	0	1	0	0	1
6	0	0	1	0	1

[5 rows x 23 columns]

```
[62]: bank.loan = bank.loan.map({"yes":1,"no":0})
```

```
[63]: bank.head()
```

```
[63]:   age  default  balance  housing  loan  day month  campaign Target \
0    58         0    2143         1     0   5  may         1      no
1    44         0     29         1     0   5  may         1      no
2    33         0      2         1     1   5  may         1      no
5    35         0    231         1     0   5  may         1      no
6    28         0    447         1     1   5  may         1      no

      blue-collar  ...  retired  self-employed  services  student  technician \
0                0  ...         0              0          0         0          0
1                0  ...         0              0          0         0          1
2                0  ...         0              0          0         0          0
5                0  ...         0              0          0         0          0
6                0  ...         0              0          0         0          0

      unemployed  married  single  secondary  tertiary
0                0         1       0          0          1
1                0         0       1          1          0
2                0         1       0          1          0
5                0         1       0          0          1
6                0         0       1          0          1
```

[5 rows x 23 columns]

```
[64]: bank.Target = bank.Target.map({"yes":1,"no":0})
```

```
[65]: bank.head()
```

```
[65]:   age  default  balance  housing  loan  day month  campaign  Target \
0    58         0    2143         1     0   5  may         1        0
1    44         0     29         1     0   5  may         1        0
2    33         0      2         1     1   5  may         1        0
5    35         0    231         1     0   5  may         1        0
6    28         0    447         1     1   5  may         1        0

      blue-collar  ...  retired  self-employed  services  student  technician \
0                0  ...         0              0          0         0          0
1                0  ...         0              0          0         0          1
2                0  ...         0              0          0         0          0
5                0  ...         0              0          0         0          0
6                0  ...         0              0          0         0          0

      unemployed  married  single  secondary  tertiary
0                0         1       0          0          1
1                0         0       1          1          0
2                0         1       0          1          0
5                0         1       0          0          1
6                0         0       1          0          1
```

[5 rows x 23 columns]

Instead of applying one hot encoding to month, we can directly encode it to integers. But, for initial analysis, let's apply one hot encoding.

```
[66]: # Perform one hot embedding and remove original month column
bank = pd.concat([bank,
                  pd.get_dummies(bank.month,drop_first=True)
                  ], axis=1).drop("month",axis=1)
bank.head()
```

```
[66]:   age  default  balance  housing  loan  day  campaign  Target  blue-collar  \
0   58         0    2143         1     0    5         1         0         0
1   44         0     29         1     0    5         1         0         0
2   33         0      2         1     1    5         1         0         0
5   35         0    231         1     0    5         1         0         0
6   28         0    447         1     1    5         1         0         0

   entrepreneur  ...  dec  feb  jan  jul  jun  mar  may  nov  oct  sep
0              0  ...   0   0   0   0   0   0   1   0   0   0
1              0  ...   0   0   0   0   0   0   1   0   0   0
2              1  ...   0   0   0   0   0   0   1   0   0   0
5              0  ...   0   0   0   0   0   0   1   0   0   0
6              0  ...   0   0   0   0   0   0   1   0   0   0
```

[5 rows x 33 columns]

```
[67]: bank.shape
```

```
[67]: (42464, 33)
```

Finally, we are left with 42464 rows and 33 features.

Before we proceed with model building, let's normalize the attribute values.

```
[68]: bank_backup = bank.copy()
```

```
[69]: Y = bank.Target
bank.drop("Target",axis=1,inplace=True)
```

```
[70]: bank = (bank - bank.min()) / (bank.max() - bank.min())
```

```
[71]: bank["Target"] = Y
```

We are now ready for model building.

1.6.5 Model Building

```
[72]: from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
```

```
from sklearn.metrics import classification_report
```

```
[73]: X = bank.drop("Target",axis=1)
      y = bank.Target
```

```
[74]: # 30% of the data will be used for testing
      test_size= 0.30
      seed = 42
      X_train, X_test, y_train, y_test= train_test_split(X, y, test_size=test_size,
      ↪random_state=seed)
```

```
[75]: y_train.shape, y_test.shape
```

```
[75]: ((29724,), (12740,))
```

There are 29.7K rows in training dataset and 12.7K rows in testing dataset.

For classification algorithms, let's start off with logistic regression, kNN, decision tree and gaussian naive bayes classifier.

```
[76]: from sklearn.linear_model import LogisticRegression
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.naive_bayes import GaussianNB
```

```
[77]: models = []
      models.append(('LR', LogisticRegression()))
      models.append(('KNN', KNeighborsClassifier()))
      models.append(('CART', DecisionTreeClassifier()))
      models.append(('NB', GaussianNB()))
```

```
[78]: results_c = []
      names_c = []

      for name, model in models:
          # define how to split off validation data ('kfold' how many folds)
          kfold = KFold(n_splits=10, random_state=seed)
          # train the model
          cv_results = cross_val_score(model, X_train, y_train, cv=kfold,
          ↪scoring='accuracy')
          results_c.append(cv_results)
          names_c.append(name)
          msg = "Train: %s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
          print(msg)
          # fit the model
          model.fit(X_train,y_train)
          # test the model
          predictions = model.predict(X_test)
          # accuracy score
          msg = "Test: %s: %f" % (name,accuracy_score(y_test, predictions))
          print(msg)
          # confusion matrix
```

```

cm = confusion_matrix(y_test, predictions)
print(cm)
# classification report
print(classification_report(y_test, predictions))

```

Train: LR: 0.884302 (0.005231)

Test: LR: 0.884615

```

[[11153  119]
 [ 1351  117]]

```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	11272
1	0.50	0.08	0.14	1468
accuracy			0.88	12740
macro avg	0.69	0.53	0.54	12740
weighted avg	0.85	0.88	0.85	12740

Train: KNN: 0.878650 (0.004175)

Test: KNN: 0.878493

```

[[10981  291]
 [ 1257  211]]

```

	precision	recall	f1-score	support
0	0.90	0.97	0.93	11272
1	0.42	0.14	0.21	1468
accuracy			0.88	12740
macro avg	0.66	0.56	0.57	12740
weighted avg	0.84	0.88	0.85	12740

Train: CART: 0.815200 (0.007050)

Test: CART: 0.822841

```

[[10098  1174]
 [ 1083   385]]

```

	precision	recall	f1-score	support
0	0.90	0.90	0.90	11272
1	0.25	0.26	0.25	1468
accuracy			0.82	12740
macro avg	0.58	0.58	0.58	12740
weighted avg	0.83	0.82	0.83	12740

Train: NB: 0.854764 (0.007146)

Test: NB: 0.858948

```

[[10521  751]]

```

```
[ 1046   422]]
      precision    recall  f1-score   support

     0       0.91      0.93      0.92     11272
     1       0.36      0.29      0.32      1468

 accuracy          0.86     12740
 macro avg       0.63      0.61      0.62     12740
weighted avg       0.85      0.86      0.85     12740
```

While the accuracy scores obtained for the test data and using KFold show that all the models perform quite well and that logistic regression performs the best, the classification report shows very low values for precision and recall for $y = 1$. This is common across all the models and is primarily because of unbalanced data. This can be fixed by upscaling the minority class ($y=1$).

```
[79]: from sklearn.utils import resample
```

```
[80]: bank.Target.value_counts()
```

```
[80]: 0    37559
      1    4905
      Name: Target, dtype: int64
```

```
[81]: # Separate majority and minority classes
df_majority = bank[bank.Target==0]
df_minority = bank[bank.Target==1]
```

```
[82]: # Upsample minority class
df_minority_upsampled = resample(df_minority,
                                replace=True,      # sample with replacement
                                n_samples=37559,    # to match majority class
                                random_state=seed)  # reproducible results
```

```
[83]: # Combine majority class with upsampled minority class
bank_upsampled = pd.concat([df_majority, df_minority_upsampled])
```

```
[84]: # Display new class counts
bank_upsampled.Target.value_counts()
```

```
[84]: 1    37559
      0    37559
      Name: Target, dtype: int64
```

```
[85]: X = bank_upsampled.drop("Target",axis=1)
      y = bank_upsampled.Target
```

```
[86]: # 30% of the data will be used for testing
test_size= 0.30
seed = 42
X_train, X_test, y_train, y_test= train_test_split(X, y, test_size=test_size,
                                                    random_state=seed)
```

```
[87]: models = []
models.append(('LR', LogisticRegression()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
models.append(('NB', GaussianNB()))

[88]: results_c = []
names_c = []

for name, model in models:
    # define how to split off validation data ('kfold' how many folds)
    kfold = KFold(n_splits=10, random_state=seed)
    # train the model
    cv_results = cross_val_score(model, X_train, y_train, cv=kfold,
    →scoring='accuracy')
    results_c.append(cv_results)
    names_c.append(name)
    msg = "Train: %s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
    # fit the model
    model.fit(X_train,y_train)
    # test the model
    predictions = model.predict(X_test)
    # accuracy score
    msg = "Test: %s: %f" % (name,accuracy_score(y_test, predictions))
    print(msg)
    # confusion matrix
    cm = confusion_matrix(y_test, predictions)
    print(cm)
    # classification report
    print(classification_report(y_test, predictions))
```

Train: LR: 0.663231 (0.004070)

Test: LR: 0.669684

[[8103 3212]

[4232 6989]]

	precision	recall	f1-score	support
0	0.66	0.72	0.69	11315
1	0.69	0.62	0.65	11221
accuracy			0.67	22536
macro avg	0.67	0.67	0.67	22536
weighted avg	0.67	0.67	0.67	22536

Train: KNN: 0.827432 (0.003217)

Test: KNN: 0.837638

[[8153 3162]

```
[ 497 10724]]
      precision    recall  f1-score   support

     0       0.94      0.72      0.82     11315
     1       0.77      0.96      0.85     11221

 accuracy      0.84     22536
 macro avg      0.86     22536
 weighted avg   0.86     22536
```

Train: CART: 0.926401 (0.003759)

Test: CART: 0.931665

```
[[ 9816 1499]
 [  41 11180]]
      precision    recall  f1-score   support

     0       1.00      0.87      0.93     11315
     1       0.88      1.00      0.94     11221

 accuracy      0.93     22536
 macro avg      0.94     22536
 weighted avg   0.94     22536
```

Train: NB: 0.617873 (0.006255)

Test: NB: 0.619542

```
[[10030 1285]
 [ 7289 3932]]
      precision    recall  f1-score   support

     0       0.58      0.89      0.70     11315
     1       0.75      0.35      0.48     11221

 accuracy      0.62     22536
 macro avg      0.67     22536
 weighted avg   0.67     22536
```

There are 2 things to note above. Even though in most cases the accuracy score has reduced, there is significant improvement in precision and recall scores for class 1. Out of the 4 models we tried, kNN and decision trees provided the best results.

1.6.6 Ensemble Techniques

Random Forest

```
[89]: from sklearn.ensemble import RandomForestClassifier
[90]: rf = RandomForestClassifier(n_estimators=10,random_state = seed)
[91]: scores = cross_val_score(rf, X_train, y_train, cv=5)
```



```
[92]: scores.mean(),scores.std()
```

```
[92]: (0.9581035763173, 0.003385948634433885)
```

Using random forest classifier, we are able to get a cross validation score of 0.9571

```
[93]: rf.fit(X_train,y_train)
```

```
[93]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=10,
                             n_jobs=None, oob_score=False, random_state=42, verbose=0,
                             warm_start=False)
```

```
[94]: # test the model
predictions = rf.predict(X_test)
# accuracy score
msg = "Test: Random Forest: %f" % (accuracy_score(y_test, predictions))
print(msg)
# confusion matrix
cm = confusion_matrix(y_test, predictions)
print(cm)
# classification report
print(classification_report(y_test, predictions))
```

Test: Random Forest: 0.967519

```
[[10654  661]
```

```
 [  71 11150]]
```

	precision	recall	f1-score	support
0	0.99	0.94	0.97	11315
1	0.94	0.99	0.97	11221
accuracy			0.97	22536
macro avg	0.97	0.97	0.97	22536
weighted avg	0.97	0.97	0.97	22536

Note how random forest classifier was able to obtain high test accuracy and high precision and recall for both classes as well.

AdaBoost Classifier

```
[95]: from sklearn.ensemble import AdaBoostClassifier
```

```
[96]: clf = AdaBoostClassifier(n_estimators=100,random_state=seed)
```

```
[97]: scores = cross_val_score(clf, X_train, y_train, cv=5)
```

```
[98]: scores.mean(),scores.std()
```

[98]: (0.6793387665127792, 0.002941049135716967)

As we can see here, AdaBoost classifier was not able to obtain high accuracy. One of the reasons behind this can be that it's made of decision stumps by default, instead of decision trees like random forest.

```
[99]: clf.fit(X_train,y_train)
```

```
[99]: AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None, learning_rate=1.0,
                        n_estimators=100, random_state=42)
```

```
[100]: # test the model
predictions = clf.predict(X_test)
# accuracy score
msg = "Test: AdaBoost Classifier: %f" % (accuracy_score(y_test, predictions))
print(msg)
# confusion matrix
cm = confusion_matrix(y_test, predictions)
print(cm)
# classification report
print(classification_report(y_test, predictions))
```

Test: AdaBoost Classifier: 0.683617

[[8416 2899]

[4231 6990]]

	precision	recall	f1-score	support
0	0.67	0.74	0.70	11315
1	0.71	0.62	0.66	11221
accuracy			0.68	22536
macro avg	0.69	0.68	0.68	22536
weighted avg	0.69	0.68	0.68	22536

Even though the accuracy score obtained was not high, the precision, recall and f1-score were around 0.6 to 0.7, which is not bad.

Gradient Boosting Classifier

```
[101]: from sklearn.ensemble import GradientBoostingClassifier
```

```
[102]: clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.
    →0,max_depth=1, random_state=seed)
```

```
[103]: scores = cross_val_score(clf, X_train, y_train, cv=5)
```

```
[104]: scores.mean(),scores.std()
```

[104]: (0.6818870723523378, 0.004003411253786577)

Similar to AdaBoost classifier, gradient boosting classifier was not able to obtain high accuracy.

```
[105]: clf.fit(X_train,y_train)
```

```
[105]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
                                learning_rate=1.0, loss='deviance', max_depth=1,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_iter_no_change=None, presort='auto',
                                random_state=42, subsample=1.0, tol=0.0001,
                                validation_fraction=0.1, verbose=0,
                                warm_start=False)
```

```
[106]: # test the model
predictions = clf.predict(X_test)
# accuracy score
msg = "Test: Gradient Boosting Classifier: %f" % (accuracy_score(y_test,
    predictions))
print(msg)
# confusion matrix
cm = confusion_matrix(y_test, predictions)
print(cm)
# classification report
print(classification_report(y_test, predictions))
```

Test: Gradient Boosting Classifier: 0.687167

[[8470 2845]

[4205 7016]]

	precision	recall	f1-score	support
0	0.67	0.75	0.71	11315
1	0.71	0.63	0.67	11221
accuracy			0.69	22536
macro avg	0.69	0.69	0.69	22536
weighted avg	0.69	0.69	0.69	22536

Similar to AdaBoost classifier, the accuracy score is not high but all precision, recall and f1-score lie in the range of 0.6-0.75

1.6.7 Model Comparison

Among all the models considered, kNN, Decision Tree and Random Forest were able to obtain high accuracy scores. Of the 3, as expected, random forest gave the highest accuracy score. On the other hand, other ensembles like gradient boosting classifier and AdaBoost classifier didn't provide any better score as compared to other classifiers like Logistic Regression and Naive Bayes.

Comparing the difference between accuracy scores for Decision Tree and Random Forest, it can be clearly observed that the ensemble technique outperforms the Decision Tree classifier. Moreover, since we didn't provide any specific parameters in Decision Tree to prevent it from overfit-

ting, it is not as good a general model as random forest classifier will be, since it's not prone to overfitting.